
Assignment 5: Fault-Tolerant Key/Value Service (7pt)

Deadline: Dec. 12, 2023; 23:59 CET

In this assignment, you will get practical experience with distributed consensus and Raft by implementing a *fault-tolerant key/value service*, which is a service a client can use to retrieve and store values using a unique key. This assignment is structured as follows:

- In the first task, you will explore how the Raft algorithm works using an interactive visualization.
- In the second task, you will implement the fault-tolerant key/value service in Python.
- In the third task, you will simulate failures to assess the fault tolerance of your implementation.
- In the fourth task, a few conceptual questions will guide you to dive deeper into Raft.

The project uses the Python *PySyncObj* package¹, which implements the Raft algorithm.

① Discovering Raft (1pt)

The website <https://raft.github.io/> includes an interactive visualization of how the Raft algorithm works in a group of five nodes. In this task, you will use this interactive visualization to perform several exercises.

Hint: Before you start working on the exercises, we also recommend going through the following visual illustration of how Raft works: <https://thesecretlivesofdata.com/raft/>.

- 1.) (0.2pt) What happens when Raft starts? Explain the process of electing a leader in the first term.
- 2.) (0.4pt) Once the leader has been elected, send a request to the leader and wait until the operation is committed by all the nodes (*note:* in the interactive visualization, a committed operation is represented with a solid border in replicated logs). Then, perform the following steps:
 - Pause the simulation, perform a new request on the leader, and take a screenshot (task1_q2_1.PNG). Stop the leader and then resume the simulation for a new election.
 - Once, there is a new leader, perform a new request, and then resume the leader from the previous term. Once the new request is committed by all servers, pause the simulation, and then take another screenshot (task1_q2_2.PNG).

How did the logs of the nodes evolve during the above exercise? Explain what happened.

- 3.) (0.4pt) Using the same visualization, stop the current leader and two additional servers. Only two nodes should now be running in your simulation. Then, perform the following steps:
 - After a few term increments, pause the simulation to take a screenshot (task1_q3_1.PNG).

¹<https://pysyncobj.readthedocs.io/en/latest/>

- Resume all nodes and resume the simulation. After the new leader election, pause the simulation and take a screenshot (task1_q3_2.PNG).

How do you explain the behavior of the system during the above exercise?

② Implementing a Fault-Tolerant Key/Value Service with Raft (2.5pts)

In this task, you will implement a key/value service that supports three types of operations:

- **Get(Key)** returns the value associated with the key. The value is a list of strings.
- **Put(Key, Value)** sets the value (a list) for the given key.
- **Append(Key, Value)** adds the value (a string) to the list of values for the given key. This operation acts like **Put(Key, [Value])** if the current value for the key is empty or undefined.

To enable fault tolerance, the service can use Raft to replicate the key/value store across a group of servers. We provide you with a Python project template based on PySyncObj², an open-source implementation of Raft in Python. To implement this task, you will then have to update the following files (see also project README):

- 1.) `kvstorage.py`, which implements an instance of the key/value service itself, and
- 2.) `kvstorage_http.py`, which implements an HTTP server that is used to perform queries on the key/value service.

Initialize three Raft servers using your key/value service implementation by using the provided initialization scripts (see README). Then, perform HTTP queries to all the three servers.

For example, to initialize the “a” key with a **Put** request, one could do the request:

```
POST http://localhost:8080/keys/a HTTP/1.1
Content-Type: application/json
```

```
{“type”: “PUT”,
“value”: [“cat”, “dog”]
}
```

Then, one can perform an **Append** request with:

```
POST http://localhost:8080/keys/a HTTP/1.1
Content-Type: application/json
```

```
{“type”: “APPEND”,
“value”: [“mouse”]
}
```

The final value can then be read with a **Get** request:

```
GET http://localhost:8080/keys/a HTTP/1.1
```

The status for node 0 can be read with:

```
GET http://localhost:8080/admin/status HTTP/1.1
```

To test your implementation:

- Perform a **Put** operation on server 0 to test your implementation.

²<https://pysyncobj.readthedocs.io/en/latest/>

- You can use the **Get** operation on server 1 with the same key to verify that the put operation was replicated and that server 1 reflects the changes.

In your **Report.md** file, indicate the replies you get from the **/admin/status** HTTP endpoint of each server. Then, answer the following questions:

- 1.) Which server is the leader? Can there be multiple leaders?
- 2.) Perform a **Put** request for the key “a” on the leader. What is the new status? What changes occurred and why (if any)?
- 3.) Perform an **Append** request for the key “a” on the leader. What is the new status? What changes occurred and why (if any)?
- 4.) Perform a **Get** request for the key “a” on the leader. What is the new status? What change (if any) happened and why?

③ Evaluation of the Key/Value Service (2pts)

In this task, you will simulate several failures to test the behavior of your implementation from **Task 2**:

- 1.) Shut down the server that acts as a leader. Report the status that you get from the servers that remain active after shutting down the leader.
- 2.) Perform a **Put** request for the key “a”. Then, restart the server from the previous point, and indicate the new status for the three servers. Indicate the result of a **Get** request for the key “a” to the previous leader.
- 3.) Has the **Put** request been replicated? Indicate which steps lead to a new election and which ones do not. Justify your answer using the statuses returned by the servers.
- 4.) Shut down two servers, including the leader — starting with the server that is not the leader. Report the status of the remaining servers and explain what happened.
- 5.) Can you perform **Get**, **Put**, or **Append** requests in this system state? Justify your answer.
- 6.) Restart the servers and note down the new status. Describe what happened.

④ Questions on Raft (1.5pt)

- 1.) (0.5pt) What is a consensus algorithm? What are they used for in the context of replicated state machines?
- 2.) (0.5pt) What are the main features of the Raft algorithm? How does Raft enable fault tolerance?
- 3.) (0.5pt) What are Byzantine failures? Can Raft handle them?

Hand-in Instructions By the deadline, you should hand in a single **zip** file via Canvas upload. The name of this file should start with **a5** and contain the last names of all team members separated by underscores (e.g., **a5_lemee_jha_ciortea.zip**). It should contain the following files:

- All answers to the assignment questions in the given **REPORT.md**; if you wish to submit your solution code via GitHub, please include a link to your GitHub repository as well

- Your code as an archive if you do not provide a GitHub link in your Report.
- The different screenshots that you have to take. (i.e., task1_q2_1.PNG, task1_q2_2.PNG, task1_q3_1.PNG, and task1_q3_2.PNG)

Across all tasks in this and the other assignments in this course, you are **required to declare** any support that you received from others and, within reasonable bounds,³ any support tools that you were using while solving the assignment.

³It is not required that you declare that you were using a text-editing software with orthographic correction; it is however required to declare if you were using any non-standard tools such as generative machine learning models (e.g., GPT).